

# Poission Smoothing Example

Doug Nychka

2025-02-18

```
library( statmod)
suppressMessages(library( fields) )
suppressMessages( library( statmod))
setwd("~/Dropbox/Home/Desktop2/PhDProjects/densityOlga")
```

## Some background

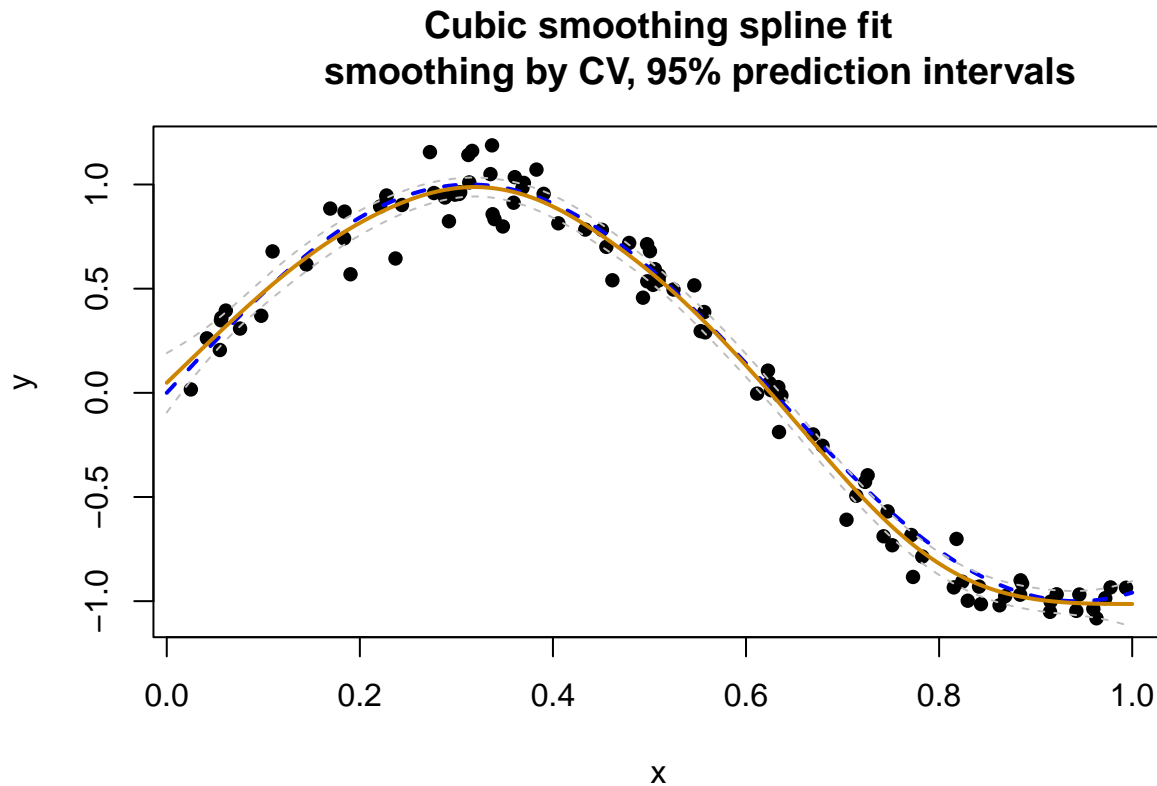
This example is an introduction to using a common curve fitting method called a spline. The function **Tps** implements this in the **fields** package and it is a likelihood based method to fit curves and surfaces. For the 1-D case in problems 1 and 2 the default method is known as a cubic smoothing spline although this is not important to do this HW.

## Using the function Tps in R.

Here is a quick 1 D example how to use **Tps**. In looking at this code observe that the **Tps** function does not take the “formula” syntax. Just give it the locations and the values for the data in that order.

```
#some simulated data
set.seed(113)
n<- 100
# random points between [0,1]
x<- sort(runif( 100))
# test function
y<- sin(5*x) + .1*rnorm( n)
# true curve we are trying to estimate
xGrid<- seq( 0,1,length.out= 300)
fTrue<- sin(5*xGrid)

# fit a spline
obj<-Tps( x, y) # amount of smoothing found by cross-validation
# take a look
plot( x,y, pch=16)
lines( xGrid, fTrue, col="blue", lty=2, lwd=2)
# predicted curve
fhat<- predict( obj,xGrid)
lines( xGrid, fhat, lwd=2,, col="orange3")
# standard errors of prediction
SE<- predictSE( obj, xGrid)
# 95% confidence envelope
lines( xGrid, fhat + 1.96* SE, col="grey", lty=2)
lines( xGrid, fhat - 1.96*SE, col="grey", lty=2)
title("Cubic smoothing spline fit
      smoothing by CV, 95% prediction intervals")
```



## More about Tps

The **Tps** object has quite a bit of information about the fit.

```
print( obj)
```

```
## Call:
## Tps(x = x, Y = y)
##
## Number of Observations:      100
## Number of parameters in the null space 2
## Parameters for fixed spatial drift    2
## Model degrees of freedom:      7.6
## Residual degrees of freedom:    92.4
## GCV estimate for tau:          0.0975
## MLE for tau:                  0.1045
## MLE for sigma:                14.04
## lambda                        0.00078
## User supplied sigma           NA
## User supplied tau^2           NA
## Summary of estimates:
##          lambda      trA      GCV      tauHat -lnLike Prof converge
## GCV      0.0007782923 7.608059 0.01028970 0.09750312 -73.55961      1
## GCV.model      NA      NA      NA      NA      NA      NA
## GCV.one      0.0007782923 7.608059 0.01028970 0.09750312      NA      1
## RMSE      NA      NA      NA      NA      NA      NA
```

## pure error	NA	NA	NA	NA	NA	NA
## REML	0.0002227867	9.967793	0.01051223	0.09728513	-77.38011	9

The error variance is estimated by “tau squared” and the estimate reported here is close to the true error standard deviation.

The

Model degrees of freedom: 7.6

Indicates that about 8 “effective” parameters are needed to represent this curve. (The number of effective parameters does not have to be a whole number.) This value, by default, is found from the data by cross-validation. But to fix it at a certain amount, for example for 3.5 degrees of freedom or 15 you can just add the argument `df`

```
obj3_5<-Tps( x, y, df= 3.5)
obj15<-Tps( x, y, df= 15)
```

```
gHat3_5<- predict(obj3_5,xGrid)
gHat15 <- predict(obj15,xGrid)
```

The returned object also includes the predicted values at the data points and the residuals using the usual format. E.g. for the fit above these are `obj$fitted.values` and `obj$residuals`.

Note that the `predict` function for the `Tps` object only returns the predicted curve, use `predictSE` to get the standard errors.


## Tornado occurence over space.

The code below reads in and then plots a version of the tornado data aggregated over space. Here the Midwest US is divided up into grid boxes (30X30) and the number of tornadoes are counted over the period 1950-2020 for each grid box. The data set is **loc**, the grid box centers, in longitude and latitude and **y** the counts in each grid box. The bubblePlot below gives a useful plot of these data.

```
load("tornadoCountsSpace.rda")


# set zero counts to grey to make plot
# easier to interpret and also just use a small number of colors
ctab<- c("grey95", tim.colors(40))

bubblePlot( loc, y,col=ctab,
             highlight=FALSE
           )
US( add=TRUE, col="magenta", lwd=2)
```



Here is a naive surface fit to these counts using **Tps**. It is not right because it does not adjust for Poisson data and also because it is not constrained to be positive. But this is an example of how to fit a spatial data set and get a quick plot the surface using **bubblePlot**. Since there are 900 locations this takes more time to fit using **Tps** than the analysis in Problems 1 and 2. The first surface plot is straight forward using the predicted values. The second evaluates on a fine grid (with the surface values in the **out\$z** component.)


```
TpsObj<- Tps( loc,y)
bubblePlot(loc, TpsObj$fitted.values,
           col=tim.colors())
US( add=TRUE, col="grey")
```



*# a more detailed plot evaluating the fit on a finer grid*

```
out<- predictSurface( TpsObj)
imagePlot( out, col=tim.colors())

US( add=TRUE, col="grey")
```



## Poisson regression

One can refit this surface fitting using a Poisson model for the counts and fit using the IRW algorithm.

The model is the number of tornados in grid box at location  $s$  is Poisson with expected value  $\exp(g(s))$  and  $g$  is assumed to be a smooth surface given by a Gaussian process – in this case a thin plate spline. Note that

in the fitting the effective degrees of freedom for the is specified and varying this will effect the smoothness of the surface.


```
nuOld<- rep( log( mean(y)), length(y))

for( k in 1:10){
  muI<- exp( nuOld)
  W<- c(muI)
  z<- nuOld + (1/muI)*( y- muI)
  # in place of WLS -- a smoothing/curve fitting step
  # note that smooting found by default method ( CV)
  tempObj<- Tps( loc,z,
                 weights=W,
                 df=60, give.warnings=FALSE)
  nuNew <- tempObj$fitted.values
  test<- sqrt( mean( (nuNew- nuOld)^2))
  cat( k, test, fill=TRUE)
  nuOld<- nuNew
}
```

```
## 1 0.7455175
## 2 0.4084566
## 3 0.2705066
## 4 0.1499811
## 5 0.03667542
## 6 0.001998531
## 7 1.179648e-05
## 8 3.166767e-05
## 9 7.953344e-07
## 10 4.08871e-13
```

```
ctab<- c("grey95", tim.colors( 64))
bubblePlot(
  loc,exp(tempObj$fitted.values),
  col=ctab,
  highlight=FALSE)

US( add=TRUE, col="magenta", lwd=2)
```



```
# a surface plot
objSur<- predictSurface( tempObj, nx=150, ny=300)
imagePlot(objSur$x, objSur$y, exp(objSur$z), col=ctab )
US( add=TRUE, col="magenta", lwd=2)
```

